

Intelligent Email Classification and Response System Using FLAN-T5

Introduction

In today's digital environment, organizations face a large volume of customer emails that include inquiries, complaints, and requests. Managing these emails effectively is crucial for maintaining customer satisfaction and smooth operational flow. However, handling emails manually can be time-consuming and prone to errors. This project presents an automated solution for classifying customer emails and generating accurate and contextually relevant responses using the FLAN-T5 language model. By leveraging Python and the Hugging Face Transformers library, this system provides a scalable tool for businesses to address customer queries with minimal human involvement.

Objective

The main goal of this project is to create an AI-powered system that classifies customer emails into predefined categories and generates empathetic, relevant, and actionable responses. By reducing manual workload, the system enhances operational efficiency without compromising customer care. FLAN-T5 enables the system to understand email content, identify customer intent, and generate responses aligned with the organization's tone and values.

System Architecture

The system consists of several modular components, each playing a key role in the email classification and response generation process. Below is an outline of its structure:

1. **Email Input Module:** Customer emails are collected into a dataset containing email IDs, subject lines, and message bodies. This dataset forms the foundation for all subsequent processing.
2. **Model Initialization:** The FLAN-T5 model and its tokenizer are retrieved from the Hugging Face Transformers library. Pre-trained on large datasets, the model is fine-tuned for tasks like email categorization and response generation.
3. **Category Framework:** A set of 15 distinct categories (e.g., "refund requests," "product questions," "technical support") is defined to classify emails, facilitating streamlined processing and response selection.

4. **Processing Pipeline:** Emails are processed sequentially through a pipeline that includes content extraction, classification, response generation, and quality assessment.
5. **Result Output:** After processing, the system exports email IDs, assigned categories, and generated responses in structured formats like CSV or JSON, enabling integration with CRM platforms or support interfaces.

Process Flow

The system follows a structured workflow to ensure accurate email classification and response generation with minimal human input.

1. **Dataset and Model Preparation:** The process begins by loading a sample dataset of customer emails, including subject lines and message bodies. The FLAN-T5 model is initialized from the Hugging Face repository to handle both classification and response generation tasks.
2. **Email Classification and Response Crafting:**
 - The email body is extracted and preprocessed for analysis.
 - FLAN-T5 assigns the email to one of the predefined categories, such as "billing inquiry" or "technical issue."
 - A response template matching the category is selected.
 - Using the template and email content, FLAN-T5 generates a concise, empathetic, and actionable reply tailored to the customer's query.
3. **Response Quality Evaluation:** To ensure optimal response quality, multiple response variants are generated using different sampling parameters (e.g., temperature, top-p sampling). These are evaluated using a custom scoring mechanism that assesses relevance, tone, and completeness. Responses that meet a quality threshold are finalized, while those that fall short are flagged for human review, ensuring consistent quality.
4. **Data Export:** Once processing is complete, results—including email IDs, categories, and responses—are exported in a structured format (CSV). This output supports further analysis or integration with existing customer service systems.

Conclusion

This project successfully demonstrates the value of using large language models like FLAN-T5 to automate and enhance customer email processing. The system effectively classifies emails into relevant categories and generates empathetic, coherent responses, significantly reducing manual workload while maintaining high communication standards.

The results show that the model handles various types of emails—such as complaints, inquiries, and requests—with a strong understanding of intent and tone. Responses are generally relevant and appropriate, supporting the system's potential for real-world applications in customer service.

While the system performs well overall, there are areas that could benefit from further refinement. For instance, some responses—though empathetic—could be more specific or complete in addressing the customer's issue. In certain cases, response templates may sound repetitive or overly generic. Introducing more dynamic response generation strategies, fine-tuning the model on a domain-specific dataset, and incorporating feedback loops for continuous learning could enhance both accuracy and personalization.